Jaime Ferrando Huertas

fhjaime96@gmail.com | www.jaimefh.com | github.com/jiwidi

Technical co-founder with deep expertise in recommendation systems, real-time ML serving, and end-to-end AI architecture. 9+ years spanning individual contributor and leadership roles across startups and enterprises. Published at NeurIPS with contributions to major open-source ML projects.

Work Experience

EcomID *Stockholm based*

Co-Founder - Head of AI *November 2023 - Present*

- Leads the AI/ML function at eComID, shaping company strategy and product roadmap as a key decision maker, defining technical direction and closely collaborating with cross-functional leadership to drive product design and user experience.
- Designed end-to-end AI architecture: data pipelines, model training infrastructure, and production deployment systems.

Shaped.ai (acquired by Whatnot) *Remote, New York*

Founding Machine Learning Engineer *Apr 2022 - Nov 2023*

- First hire. Grew from engineering into system architecture and product, working directly with customers to set the product direction.
- Built the core recommendation and ranking model library: collaborative filtering (ALS, SVD, LightFM), sequential models (GRU4Rec, BERT4Rec), two-tower models and learning-to-rank (LightGBM, XGBoost), powering real-time recommendations for 50+ customers.
- Built vector retrieval (FAISS, Annoy), candidate generation, session-based ranking and exploration (Thompson sampling, epsilon-greedy) behind low-latency ranking APIs, plus semantic search on top of the recommendation model representations.
- Designed the real-time feature store and data ingestion pipelines (BigQuery, PostgreSQL, MySQL, Airflow, Redis), merging batch and streaming to minimize train-inference skew.
- Built the training platform on SageMaker: large-scale hyperparameter sweeps, automated model benchmarking and machine sizing from memory estimates to keep compute cost down.
- Owned the AWS infrastructure end to end: CDK infrastructure as code, EKS, Lambda, CI/CD, end-to-end tests and production rollouts.

H&M *Stockholm, Sweden*

Senior Machine Learning Engineer *Feb 2020 - Apr 2022*

- Organized H&M's first Kaggle competition (2,900+ teams, $50,000 prize fund) and released the open H&M recommendation dataset for it.
- Engineered neural-based recommendation systems, significantly boosting customer engagement, revenue and online interaction.
- Designed and implemented infrastructure on GCP for real-time recommendation model serving. Reduced ingestion time of user events from 16 hours to 3 seconds.
- Steered A/B tests and experimentation, and set up feedback loops for internal dashboards.

Sandvik *Stockholm, Sweden*

Data Scientist *Oct 2018 - Feb 2020*

- Developed anomaly detection models for multivariate data in high-performance industrial processes.
- Deployed and maintained inference servers on GCP.

Polystar group *Stockholm, Sweden*

Data Scientist *Oct 2017 - Jul 2018*

- Developed anomaly detection models for telecommunications networks.

Ahora freeware *Valencia, Spain*

Software Engineer Intern *Jun 2016 - Oct 2017*

- Designed a cloud-based library integrating with existing CRM for file storage across Dropbox, Office365, and Google Drive.

Education

Master's degree in Artificial Intelligence, Pattern Recognition and Digital Imaging *Valencia, Spain*

UPV, Polytechnic University of Valencia *Aug 2020 - Jun 2021*

- Master's thesis (cum laude) on neural recommender systems at H&M. The proposed models replaced the ones in production.

Bachelor in Computer Science *Valencia, Spain*

UPV, Polytechnic University of Valencia *Aug 2014 - Jun 2018*

- Class representative.
- Selected for the high-performance group taught in English (30 of 450 students).
- Exchange program at KTH Royal Institute of Technology.

Publications

Explaining Preferences with Shapley Values *NeurIPS 2022*
Robert Hu, Siu Lun Chau, Jaime Ferrando Huertas, Dino Sejdinovic

Projects and Open Source

jiwo
Small, low-latency decision models: one forward pass returns a probability per option, with no text generation. At training time, jiwo-0.8b ranked 1st of 27 sub-1B models on the Decision Index. Nearly 300 GitHub stars.

HowTo
Rust CLI that turns plain English into a shell command with a local model. The models come from sandbox-verified distillation (howto-modelling). LoRA raised a 2B model from 27% to 74.5% pass@1 on held-out tasks.

lens-database.com
Archive of photography lenses. A pipeline scrapes listings from multiple sources, deduplicates them into one canonical catalog and enriches them with an LLM. 10k monthly users.

scikit-learn
Several improvements to the Partial Least Squares (PLS) implementation and helper methods.

PyTorch Audio
Added new baseline datasets (TED-LIUM 1, 2 and 3) and code to facilitate reproducible experiments on one of the most used ASR benchmarks.

Time series with Python
A use-case focused tutorial for time series forecasting with Python. Nearly 700 GitHub stars.

Behavior-Sequence-Transformer-PyTorch
Open-source implementation of Alibaba's Behavior Sequence Transformer, which combines user behavior sequences and context features for recommendation. 170+ GitHub stars.

LAS-PyTorch
Open-source implementation of the Listen, Attend and Spell model for end-to-end speech recognition.

Achievements

2022 **Organizer,** H&M Personalized Fashion Recommendations *Kaggle*
2018 **2nd,** UniversityHack 2018, national inter-university data science competition *Spain*

Skills

Programming	Python, Rust, SQL, Java, C++, JavaScript
ML/AI	PyTorch, scikit-learn, LightGBM, XGBoost, TensorFlow, FAISS, Hugging Face, LangChain, vLLM
Cloud & Infrastructure	AWS (SageMaker, EKS, Lambda, S3, CDK), GCP (BigQuery, Vertex AI, Dataflow, Pub/Sub), Azure
	Docker, Kubernetes, Terraform, Airflow, Jenkins, CircleCI, MLflow, Weights & Biases
Data	PostgreSQL, Redis, ClickHouse, MongoDB, Cassandra, Elasticsearch, Spark, Databricks, Kafka, Feast
Miscellaneous	Linux, Shell (Bash/Zsh), Git, System Design, A/B Testing, Real-time ML Serving

References available upon request.